

CIND820_initial_analysis

June 25, 2023

```
[ ]: import pandas as pd

df = pd.read_csv("../data/raw/Crash_Reporting_Drivers_Data.csv")
df.head()
```

```
C:\Users\Zhaokai\AppData\Local\Temp\ipykernel_19140\4144606215.py:3:
DtypeWarning: Columns (1) have mixed types. Specify dtype option on import or
set low_memory=False.
```

```
df = pd.read_csv("../data/raw/Crash_Reporting_Drivers_Data.csv")
```

```
[ ]: Report Number Local Case Number Agency Name \
0 DD5620004G 190046109 Rockville Police Departme
1 MCP29620057 200023865 Montgomery County Police
2 MCP2651004P 210035267 Montgomery County Police
3 MCP3050004K 200016465 Montgomery County Police
4 MCP2641001J 200016526 Montgomery County Police
```

```
ACRS Report Type Crash Date/Time Route Type Road Name \
0 Property Damage Crash 09/26/2019 07:20:00 AM NaN NaN
1 Property Damage Crash 06/18/2020 02:00:00 AM County DAIRYMAID DR
2 Property Damage Crash 09/08/2021 12:30:00 PM NaN NaN
3 Property Damage Crash 04/19/2020 03:39:00 PM County RUSSETT RD
4 Injury Crash 04/20/2020 09:15:00 AM County ARCHDALE RD
```

```
Cross-Street Type Cross-Street Name Off-Road Description \
0 NaN NaN PARKING LOT
1 County METZ DR NaN
2 NaN NaN PARKING LOT OF 10109 COLESVILLE RD
3 County ARCTIC AVE NaN
4 County GUNNERS BRANCH RD NaN
```

```
... Speed Limit Driverless Vehicle Parked Vehicle Vehicle Year \
0 ... 15 No No 2017
1 ... 35 No No 2020
2 ... 0 No No 2010
3 ... 25 No No 2004
4 ... 25 No No 2006
```

	Vehicle Make	Vehicle Model	Equipment Problems	Latitude	Longitude	\
0	THOMAS	BUS	NO MISUSE	39.103518	-77.157669	
1	UNK	UNK	UNKNOWN	39.154847	-77.271245	
2	TOYOTA	SUV	UNKNOWN	39.020303	-77.011436	
3	DODGE	GRAND CARAVAN	NO MISUSE	39.080062	-77.097845	
4	HONDA	CR-V	UNKNOWN	39.175230	-77.241090	

	Location
0	(39.10351817, -77.15766933)
1	(39.15484667, -77.271245)
2	(39.02030267, -77.01143583)
3	(39.08006167, -77.097845)
4	(39.17523, -77.24109)

[5 rows x 43 columns]

```
[ ]: # number of rows, number of columns for dataset
nrow,ncol=df.shape
nrow,ncol
```

```
[ ]: (159357, 43)
```

```
[ ]: # We have following data types of variables float64(2), int64(2), object(39).
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 159357 entries, 0 to 159356
```

```
Data columns (total 43 columns):
```

#	Column	Non-Null Count	Dtype
0	Report Number	159357 non-null	object
1	Local Case Number	159357 non-null	object
2	Agency Name	159357 non-null	object
3	ACRS Report Type	159357 non-null	object
4	Crash Date/Time	159357 non-null	object
5	Route Type	143710 non-null	object
6	Road Name	144586 non-null	object
7	Cross-Street Type	143684 non-null	object
8	Cross-Street Name	144577 non-null	object
9	Off-Road Description	14769 non-null	object
10	Municipality	17700 non-null	object
11	Related Non-Motorist	5019 non-null	object
12	Collision Type	158820 non-null	object
13	Weather	146879 non-null	object
14	Surface Condition	140664 non-null	object
15	Light	158035 non-null	object
16	Traffic Control	135447 non-null	object

17	Driver Substance Abuse	130386	non-null	object
18	Non-Motorist Substance Abuse	3976	non-null	object
19	Person ID	159357	non-null	object
20	Driver At Fault	159357	non-null	object
21	Injury Severity	159357	non-null	object
22	Circumstance	29401	non-null	object
23	Driver Distracted By	159357	non-null	object
24	Drivers License State	150465	non-null	object
25	Vehicle ID	159357	non-null	object
26	Vehicle Damage Extent	159077	non-null	object
27	Vehicle First Impact Location	159201	non-null	object
28	Vehicle Second Impact Location	159101	non-null	object
29	Vehicle Body Type	156923	non-null	object
30	Vehicle Movement	159018	non-null	object
31	Vehicle Continuing Dir	156880	non-null	object
32	Vehicle Going Dir	156880	non-null	object
33	Speed Limit	159357	non-null	int64
34	Driverless Vehicle	159357	non-null	object
35	Parked Vehicle	159357	non-null	object
36	Vehicle Year	159357	non-null	int64
37	Vehicle Make	159335	non-null	object
38	Vehicle Model	159295	non-null	object
39	Equipment Problems	127136	non-null	object
40	Latitude	159357	non-null	float64
41	Longitude	159357	non-null	float64
42	Location	159357	non-null	object

dtypes: float64(2), int64(2), object(39)
memory usage: 52.3+ MB

```
[ ]: # according to the output of last code section df.info(), we will remove
      ↪ columns with small amount of non-null count values
new_df=df.drop(["Off-Road Description","Municipality","Related_
      ↪Non-Motorist","Non-Motorist Substance Abuse","Circumstance"],axis=1)

# we will also remove "Driverless Vehicle", according to its values_
      ↪distribution below, we can tell this column is meaningless.
# df["Driverless Vehicle"].value_counts()
# No          158668
# Unknown      689
# Name: Driverless Vehicle, dtype: int64
new_df=new_df.drop(["Driverless Vehicle"],axis=1)

# we also remove columns that are not related to this study.
new_df=new_df.drop(["Road Name","Cross-Street Type","Cross-Street Name","Report_
      ↪Number","Local Case Number","Agency Name",
      "Person ID", "Drivers License State","Vehicle_
      ↪ID","Latitude","Longitude"],axis=1)
```

```

# During the data cleaning stage, the researcher found out most values from the
↳ columns related to the vehicle category are
# corrupted and contaminated. As a consequence, the researcher had to give up
↳ the factors under the vehicle categories.
new_df=new_df.drop(["Vehicle Year","Vehicle Make","Vehicle Model","Vehicle
↳ Going Dir","Vehicle Continuing Dir",
                    "Vehicle Movement","Vehicle Body Type","Vehicle Second
↳ Impact Location","Vehicle First Impact Location",
                    "Parked Vehicle","Vehicle Damage Extent","Equipment
↳ Problems","Location"],axis=1)

# Considering the ACRS Report Type is more similar to another target attribute,
↳ this study's class attribute is injury severity.
# As a result, this column needs to be removed.
print(new_df["ACRS Report Type"].value_counts())
new_df=new_df.drop(["ACRS Report Type"],axis=1)
new_df.info()

```

```

Property Damage Crash      101048
Injury Crash                57913
Fatal Crash                 396
Name: ACRS Report Type, dtype: int64
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 159357 entries, 0 to 159356
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Crash Date/Time                      159357 non-null object
1   Route Type                          143710 non-null object
2   Collision Type                      158820 non-null object
3   Weather                             146879 non-null object
4   Surface Condition                   140664 non-null object
5   Light                              158035 non-null object
6   Traffic Control                    135447 non-null object
7   Driver Substance Abuse              130386 non-null object
8   Driver At Fault                    159357 non-null object
9   Injury Severity                    159357 non-null object
10  Driver Distracted By                159357 non-null object
11  Speed Limit                        159357 non-null int64
dtypes: int64(1), object(11)
memory usage: 14.6+ MB

```

```

[ ]: new_df["Crash Date/Time"].value_counts()
# There is no missing value for this field.
# However, the Crash Date/Time can not be used directly for study,
# because the length is too big to manage

```

```
# So, it is a more reasonable way to convert this column into two new columns
# 1. day of the week
# 2. time of the day
```

```
[ ]: 12/10/2018 06:10:00 PM    11
      06/09/2020 06:53:00 PM    10
      03/03/2017 06:00:00 AM    10
      03/28/2019 09:30:00 AM    10
      05/17/2017 03:36:00 PM     9
      ..
      09/10/2015 10:30:00 PM     1
      06/30/2018 02:02:00 AM     1
      08/23/2018 08:32:00 PM     1
      02/20/2023 02:05:00 PM     1
      05/09/2023 02:56:00 AM     1
      Name: Crash Date/Time, Length: 87470, dtype: int64
```

```
[ ]: # Adding column Day of Week
      from datetime import datetime
      def convert_to_day_of_week(date_format:str,date_string:str):
          date_object = datetime.strptime(date_string,date_format)
          day_of_week = date_object.weekday()
          day_name_list = ["Monday", "Tuesday", "Wednesday", "Thursday", "Friday",
                           ↪ "Saturday", "Sunday"]
          day_name = day_name_list[day_of_week]
          return day_name

      def apply_conversion(datetime):
          return convert_to_day_of_week("%m/%d/%Y %I:%M:%S %p", datetime)

      new_df["Day of Week"]=new_df["Crash Date/Time"].apply(apply_conversion)
      new_df["Day of Week"].value_counts()
```

```
[ ]: Friday      25701
      Tuesday     24992
      Thursday    24831
      Wednesday   24554
      Monday      23034
      Saturday    19790
      Sunday      16455
      Name: Day of Week, dtype: int64
```

```
[ ]: # Adding column Time of Day
      from datetime import datetime
      def convert_to_time_of_day(date_format:str,date_string:str):
          date_object = datetime.strptime(date_string,date_format)
          hour = date_object.hour
```

```

if 0<= hour <6:
    return "dawn"
elif 6 <= hour < 12:
    return "morning"
elif 12<= hour <18:
    return "afternoon"
else:
    return "evening"

def apply_conversion(datetime):
    return convert_to_time_of_day("%m/%d/%Y %I:%M:%S %p", datetime)

new_df["Time of Day"]=new_df["Crash Date/Time"].apply(apply_conversion)
new_df["Time of Day"].value_counts()

```

```

[ ]: afternoon    65690
     morning      45568
     evening      37233
     dawn         10866
     Name: Time of Day, dtype: int64

```

```

[ ]: # Now the Time of Day and Day of Week Columns are generated.
     # Therefore, the Crash Date/Time can be removed.
     new_df = new_df.drop(["Crash Date/Time"],axis=1)
     new_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 159357 entries, 0 to 159356
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Route Type            143710 non-null object
 1   Collision Type        158820 non-null object
 2   Weather               146879 non-null object
 3   Surface Condition     140664 non-null object
 4   Light                 158035 non-null object
 5   Traffic Control       135447 non-null object
 6   Driver Substance Abuse 130386 non-null object
 7   Driver At Fault       159357 non-null object
 8   Injury Severity       159357 non-null object
 9   Driver Distracted By  159357 non-null object
10   Speed Limit           159357 non-null int64
11   Day of Week           159357 non-null object
12   Time of Day           159357 non-null object
dtypes: int64(1), object(12)
memory usage: 15.8+ MB

```

```
[ ]: print(new_df["Route Type"].value_counts())
# Based on the output of the value counts,
# There are three types of rows that needed to be removed
# Due to the ambiguous meaning of the value
# 1. other public roadway
# 2. government
# 3. unknown

new_df=new_df.loc[(new_df["Route Type"]!="Government")
                  &(new_df["Route Type"]!="Other Public Roadway")
                  &(new_df["Route Type"]!="Unknown")]
print("-----After removal-----")
print(new_df["Route Type"].value_counts())
```

Maryland (State)	71500
County	51339
Municipality	8713
US (State)	6926
Interstate (State)	2985
Other Public Roadway	1053
Government	612
Ramp	523
Service Road	40
Unknown	19

Name: Route Type, dtype: int64

-----After removal-----

Maryland (State)	71500
County	51339
Municipality	8713
US (State)	6926
Interstate (State)	2985
Ramp	523
Service Road	40

Name: Route Type, dtype: int64

```
[ ]: print(new_df["Collision Type"].value_counts())
# Based on the output of the value counts,
# The rows contains unknown in collision type column will be removed
new_df=new_df.loc[(new_df["Collision Type"]!="UNKNOWN")]
# The reason why we did not impute the "OTHER" value
# It could be a type different from all these types of accidents
# Other here is meaningful
print("-----After removal-----")
print(new_df["Collision Type"].value_counts())
```

SAME DIR REAR END	51626
STRAIGHT MOVEMENT ANGLE	27621
OTHER	17384

SAME DIRECTION SIDESWIPE	14813
SINGLE VEHICLE	14437
HEAD ON LEFT TURN	11723
SAME DIRECTION RIGHT TURN	3479
SAME DIRECTION LEFT TURN	3425
HEAD ON	3417
OPPOSITE DIRECTION SIDESWIPE	2618
ANGLE MEETS LEFT TURN	1870
ANGLE MEETS RIGHT TURN	1114
SAME DIR REND LEFT TURN	685
SAME DIR REND RIGHT TURN	683
UNKNOWN	667
SAME DIR BOTH LEFT TURN	648
ANGLE MEETS LEFT HEAD ON	639
OPPOSITE DIR BOTH LEFT TURN	294

Name: Collision Type, dtype: int64

-----After removal-----

SAME DIR REAR END	51626
STRAIGHT MOVEMENT ANGLE	27621
OTHER	17384
SAME DIRECTION SIDESWIPE	14813
SINGLE VEHICLE	14437
HEAD ON LEFT TURN	11723
SAME DIRECTION RIGHT TURN	3479
SAME DIRECTION LEFT TURN	3425
HEAD ON	3417
OPPOSITE DIRECTION SIDESWIPE	2618
ANGLE MEETS LEFT TURN	1870
ANGLE MEETS RIGHT TURN	1114
SAME DIR REND LEFT TURN	685
SAME DIR REND RIGHT TURN	683
SAME DIR BOTH LEFT TURN	648
ANGLE MEETS LEFT HEAD ON	639
OPPOSITE DIR BOTH LEFT TURN	294

Name: Collision Type, dtype: int64

```
[ ]: print(new_df["Weather"].value_counts())
new_df=new_df.loc[(new_df["Weather"]!="UNKNOWN")&
                  (new_df["Weather"]!="OTHER")]
print("-----After removal-----")
print(new_df["Weather"].value_counts())
```

CLEAR	106110
RAINING	18884
CLOUDY	15991
SNOW	1401
FOGGY	624
UNKNOWN	517

WINTRY MIX	381
OTHER	337
SLEET	209
SEVERE WINDS	153
BLOWING SNOW	111
BLOWING SAND, SOIL, DIRT	13

Name: Weather, dtype: int64

-----After removal-----

CLEAR	106110
RAINING	18884
CLOUDY	15991
SNOW	1401
FOGGY	624
WINTRY MIX	381
SLEET	209
SEVERE WINDS	153
BLOWING SNOW	111
BLOWING SAND, SOIL, DIRT	13

Name: Weather, dtype: int64

```
[ ]: print(new_df["Surface Condition"].value_counts())
new_df=new_df.loc[(new_df["Surface Condition"]!="UNKNOWN")&
                  (new_df["Surface Condition"]!="OTHER")]
print("-----After removal-----")
print(new_df["Surface Condition"].value_counts())
```

DRY	109664
WET	25982
ICE	936
SNOW	917
UNKNOWN	211
SLUSH	192
OTHER	155
MUD, DIRT, GRAVEL	45
WATER(STANDING/MOVING)	37
OIL	24
SAND	4

Name: Surface Condition, dtype: int64

-----After removal-----

DRY	109664
WET	25982
ICE	936
SNOW	917
SLUSH	192
MUD, DIRT, GRAVEL	45
WATER(STANDING/MOVING)	37
OIL	24
SAND	4

Name: Surface Condition, dtype: int64

```
[ ]: print(new_df["Light"].value_counts())
new_df=new_df.loc[(new_df["Light"]!="UNKNOWN")&
                  (new_df["Light"]!="OTHER")&
                  (new_df["Light"]!="DARK -- UNKNOWN LIGHTING")]
print("-----After removal-----")
new_df["Light"].value_counts()
```

DAYLIGHT	105499
DARK LIGHTS ON	35966
DARK NO LIGHTS	4486
DUSK	3604
DAWN	3090

Name: Light, dtype: int64

-----After removal-----

```
[ ]: DAYLIGHT          105499
      DARK LIGHTS ON   35966
      DARK NO LIGHTS   4486
      DUSK              3604
      DAWN              3090
Name: Light, dtype: int64
```

```
[ ]: print(new_df["Traffic Control"].value_counts())
new_df=new_df.loc[(new_df["Traffic Control"]!="UNKNOWN")&
                  (new_df["Traffic Control"]!="OTHER")
                  ]
print("-----After removal-----")
new_df["Traffic Control"].value_counts()
```

NO CONTROLS	60897
TRAFFIC SIGNAL	53135
STOP SIGN	11159
FLASHING TRAFFIC SIGNAL	1890
OTHER	1793
YIELD SIGN	1620
PERSON	256
WARNING SIGN	146
UNKNOWN	120

RAILWAY CROSSING DEVICE	42
SCHOOL ZONE SIGN DEVICE	14

Name: Traffic Control, dtype: int64

-----After removal-----

```
[ ]: NO CONTROLS          60897
      TRAFFIC SIGNAL       53135
      STOP SIGN           11159
```

```
FLASHING TRAFFIC SIGNAL      1890
YIELD SIGN                  1620
PERSON                      256
WARNING SIGN                146
RAILWAY CROSSING DEVICE      42
SCHOOL ZONE SIGN DEVICE      14
Name: Traffic Control, dtype: int64
```

```
[ ]: print(new_df["Driver Substance Abuse"].value_counts())
new_df=new_df.loc[(new_df["Driver Substance Abuse"]!="OTHER")]
print("-----After removal-----")
print(new_df["Driver Substance Abuse"].value_counts())

# if it is not non detected, unknown, or null value
# the value will be converted to detected.
def convert_substance_abuse_levels(src):
    if(src == "NONE DETECTED" or src == "UNKNOWN" or src == None):
        return src
    else:
        return "DETECTED"

new_df["Driver Substance Abuse"]=new_df["Driver Substance Abuse"].
    ↪map(convert_substance_abuse_levels)
print("-----After Converting-----")
print(new_df["Driver Substance Abuse"].value_counts())
# replace unknown with null values and then, forward fill imputation for null
↪values.
new_df["Driver Substance Abuse"] = new_df["Driver Substance Abuse"].
    ↪replace("UNKNOWN",None)
new_df["Driver Substance Abuse"] = new_df["Driver Substance Abuse"].
    ↪fillna(method="ffill")
print("-----After imputation-----")
print(new_df["Driver Substance Abuse"] .value_counts())
```

```
NONE DETECTED              109687
UNKNOWN                    9467
ALCOHOL PRESENT            3609
ALCOHOL CONTRIBUTED        1269
ILLEGAL DRUG PRESENT       238
MEDICATION PRESENT         108
ILLEGAL DRUG CONTRIBUTED   89
COMBINED SUBSTANCE PRESENT 82
MEDICATION CONTRIBUTED     57
OTHER                      57
COMBINATION CONTRIBUTED    43
Name: Driver Substance Abuse, dtype: int64
-----After removal-----
```

```

NONE DETECTED          109687
UNKNOWN                9467
ALCOHOL PRESENT        3609
ALCOHOL CONTRIBUTED    1269
ILLEGAL DRUG PRESENT   238
MEDICATION PRESENT     108
ILLEGAL DRUG CONTRIBUTED 89
COMBINED SUBSTANCE PRESENT 82
MEDICATION CONTRIBUTED  57
COMBINATION CONTRIBUTED 43
Name: Driver Substance Abuse, dtype: int64
-----After Converting-----
NONE DETECTED          109687
DETECTED               32796
UNKNOWN               9467
Name: Driver Substance Abuse, dtype: int64
-----After imputation-----
NONE DETECTED          116894
DETECTED               35056
Name: Driver Substance Abuse, dtype: int64

```

```

[ ]: print(new_df["Driver At Fault"].value_counts())
new_df=new_df.loc[(new_df["Driver At Fault"]!="Unknown")]
print("-----After removal-----")
new_df["Driver At Fault"].value_counts()

```

```

Yes          81064
No           67055
Unknown       3831
Name: Driver At Fault, dtype: int64
-----After removal-----

```

```

[ ]: Yes      81064
     No       67055
     Name: Driver At Fault, dtype: int64

```

```

[ ]: # same process as handling "Driver Substance Abuse" is performed here.
print(new_df["Driver Distracted By"].value_counts())

def convert_distracted_levels(src):
    if(src == "NOT DISTRACTED" or src == "UNKNOWN" or src == None):
        return src
    else:
        return "DISTRACTED"

new_df["Driver Distracted By"]=new_df["Driver Distracted By"].
    ↪map(convert_distracted_levels)

```

```

print("-----After removal-----")
print(new_df["Driver Distracted By"].value_counts())

# Imputes unknown field. Using forward imputation.
print("-----After imputation-----")
new_df["Driver Distracted By"] = new_df["Driver Distracted By"].
    ↪replace("UNKNOWN",None)
new_df["Driver Distracted By"] = new_df["Driver Distracted By"].
    ↪fillna(method="ffill")
new_df["Driver Distracted By"] .value_counts()

```

```

NOT DISTRACTED          93773
UNKNOWN                 25964
LOOKED BUT DID NOT SEE  18661
INATTENTIVE OR LOST IN THOUGHT  3790
OTHER DISTRACTION       2825
DISTRACTED BY OUTSIDE PERSON OBJECT OR EVENT  855
BY OTHER OCCUPANTS      381
OTHER CELLULAR PHONE RELATED  332
OTHER ELECTRONIC DEVICE (NAVIGATIONAL PALM PILOT)  301
TALKING OR LISTENING TO CELLULAR PHONE  261
NO DRIVER PRESENT       221
BY MOVING OBJECT IN VEHICLE  191
EATING OR DRINKING      177
ADJUSTING AUDIO AND OR CLIMATE CONTROLS  123
USING OTHER DEVICE CONTROLS INTEGRAL TO VEHICLE  80
TEXTING FROM A CELLULAR PHONE  59
USING DEVICE OBJECT BROUGHT INTO VEHICLE  56
DIALING CELLULAR PHONE   44
SMOKING RELATED          25

```

Name: Driver Distracted By, dtype: int64

-----After removal-----

```

NOT DISTRACTED    93773
DISTRACTED        28382
UNKNOWN           25964

```

Name: Driver Distracted By, dtype: int64

-----After imputation-----

```

[ ]: NOT DISTRACTED    113632
     DISTRACTED        34487
     Name: Driver Distracted By, dtype: int64

```

```

[ ]: # summary statistics for numeric attributes
     new_df.describe()
     # Although "Speed Limit" attribute is a numeric attribute
     # The data is organized as categorical.
     # We will reduce levels from 15 to 5.

```

```

print(new_df["Speed Limit"].value_counts())
levels = ["0-10", "15-25", "30-40", "45-55", "60-70"]
bins = [0, 10, 25, 40, 55, 70]
new_df["Speed Limit"] = pd.cut(new_df["Speed Limit"], bins= bins, labels=levels)
new_df["Speed Limit"] = new_df["Speed Limit"].astype("category")
print("-----After convert-----")
print(new_df["Speed Limit"].value_counts())

```

```

35    43337
40    29880
30    20084
25    19985
45    11162
15     5056
50     4184
0      3930
55     3534
5      3424
10     2499
20      925
60       62
65       52
70        5
Name: Speed Limit, dtype: int64
-----After convert-----
30-40    93301
15-25    25966
45-55    18880
0-10     5923
60-70     119
Name: Speed Limit, dtype: int64

```

```

[ ]: # hint: use this before you print large amount of data
pd.set_option("display.max_rows", 500)
pd.set_option("display.max_columns", 500)
pd.set_option("display.width", 1000)

```

```

[ ]: # remove all rows that contain null values.
new_df = new_df.dropna(how="any")
new_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 105379 entries, 3 to 159356
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Route Type            105379 non-null object

```

```

1 Collision Type      105379 non-null object
2 Weather            105379 non-null object
3 Surface Condition  105379 non-null object
4 Light              105379 non-null object
5 Traffic Control    105379 non-null object
6 Driver Substance Abuse 105379 non-null object
7 Driver At Fault    105379 non-null object
8 Injury Severity    105379 non-null object
9 Driver Distracted By 105379 non-null object
10 Speed Limit        105379 non-null category
11 Day of Week        105379 non-null object
12 Time of Day        105379 non-null object
dtypes: category(1), object(12)
memory usage: 10.6+ MB

```

```

[ ]: # Count the occurrences of each unique value in the class attribute "Injury
↳Severity"
new_df["Injury Severity"].value_counts()

```

```

[ ]: NO APPARENT INJURY      84376
POSSIBLE INJURY             11873
SUSPECTED MINOR INJURY      8108
SUSPECTED SERIOUS INJURY    956
FATAL INJURY                 66
Name: Injury Severity, dtype: int64

```

```

[ ]: # according to the output of last code section, 130578(NO APPARENT INJURY)/
↳159357(total rows) = 81.9%
# we can tell that the multi-class problem has imbalanced class distributions.
↳reducing number of classes could help address the issue. we will deal with
# imbalanced class distributions later.
# A large number of classes can increase the complexity of the classification
↳problem, by reducing the number of classes, we can simplify the problem
# and make it more manageable for the model to learn the distinguishing
↳features, which may lead to better performance.

# for target variable "Injury Severity", convert 5 classes to 3 classes
new_df["Injury Severity"]=new_df["Injury Severity"].map({"NO APPARENT INJURY":
↳"No Injury", "POSSIBLE INJURY": "Minor Injury", "SUSPECTED MINOR INJURY": "Minor
↳Injury", "SUSPECTED SERIOUS INJURY": "Serious Injury",
"Fatal Injury": "Serious Injury"})
new_df["Injury Severity"].value_counts()

```

```

[ ]: No Injury      84376
Minor Injury       19981
Serious Injury     1022
Name: Injury Severity, dtype: int64

```

```
[ ]: # convert each column data type from object to category, to improve efficiency
      ↪during data modeling stage.
new_df = new_df.apply(lambda x:x.astype('category'))
new_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 105379 entries, 3 to 159356
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Route Type                            105379 non-null category
1   Collision Type                        105379 non-null category
2   Weather                              105379 non-null category
3   Surface Condition                    105379 non-null category
4   Light                                105379 non-null category
5   Traffic Control                      105379 non-null category
6   Driver Substance Abuse               105379 non-null category
7   Driver At Fault                     105379 non-null category
8   Injury Severity                     105379 non-null category
9   Driver Distracted By                105379 non-null category
10  Speed Limit                          105379 non-null category
11  Day of Week                         105379 non-null category
12  Time of Day                         105379 non-null category
dtypes: category(13)
memory usage: 2.1 MB
```

```
[ ]: # summary statistics for all columns
new_df.describe()
```

```
[ ]:
count      Route Type      Collision Type  Weather  Surface Condition  \
unique              7          17          10              9
top    Maryland (State)  SAME DIR REAR END    CLEAR              DRY
freq              54209          38119    76845              83103

count      Light Traffic Control  Driver Substance Abuse  Driver At Fault  \
unique              5              9              2              2
top    DAYLIGHT  TRAFFIC SIGNAL          NONE DETECTED              Yes
freq              73360          46384              84640              55157

count      Injury Severity  Driver Distracted By  Speed Limit  Day of Week  \
unique              3              2              5              7
top      No Injury      NOT DISTRACTED      30-40      Friday
freq              84376          81402          73969          17081
```


	Time of Day
count	105379
unique	4
top	afternoon
freq	43517

```
[ ]: new_df.to_csv("cleaned_data.csv", index=False)
```

```
[ ]: # Univariate visualization for all attributes

import matplotlib.pyplot as plt

col_summary = {}
for column in new_df.columns:
    col_summary[column] = new_df[column].value_counts()

# Print the summary statistics
for column, summary in col_summary.items():
    print(f"Summary statistics for {column}:")
    print(summary)
    print()

# Visualize the summary statistics using a bar plot
for column, summary in col_summary.items():
    summary.plot(kind='bar')
    plt.title(f"Frequency of Categories in {column}")
    plt.xlabel("Categories")
    plt.ylabel("Frequency")
    plt.show()
```

Summary statistics for Route Type:

Maryland (State)	54209
County	37754
Municipality	6272
US (State)	4846
Interstate (State)	1883
Ramp	391
Service Road	24

Name: Route Type, dtype: int64

Summary statistics for Collision Type:

SAME DIR REAR END	38119
STRAIGHT MOVEMENT ANGLE	21153
HEAD ON LEFT TURN	9488
SAME DIRECTION SIDESWIPE	9378
SINGLE VEHICLE	7760
OTHER	7093

SAME DIRECTION RIGHT TURN	2320
SAME DIRECTION LEFT TURN	2189
HEAD ON	2095
ANGLE MEETS LEFT TURN	1461
OPPOSITE DIRECTION SIDESWIPE	1394
ANGLE MEETS RIGHT TURN	852
SAME DIR REND RIGHT TURN	509
ANGLE MEETS LEFT HEAD ON	478
SAME DIR REND LEFT TURN	472
SAME DIR BOTH LEFT TURN	445
OPPOSITE DIR BOTH LEFT TURN	173

Name: Collision Type, dtype: int64

Summary statistics for Weather:

CLEAR	76845
RAINING	14512
CLOUDY	11951
SNOW	992
FOGGY	475
WINTY MIX	270
SLEET	137
SEVERE WINDS	112
BLOWING SNOW	78
BLOWING SAND, SOIL, DIRT	7

Name: Weather, dtype: int64

Summary statistics for Surface Condition:

DRY	83103
WET	20743
SNOW	681
ICE	643
SLUSH	134
WATER(STANDING/MOVING)	29
MUD, DIRT, GRAVEL	27
OIL	18
SAND	1

Name: Surface Condition, dtype: int64

Summary statistics for Light:

DAYLIGHT	73360
DARK LIGHTS ON	24422
DARK NO LIGHTS	3001
DUSK	2381
DAWN	2215

Name: Light, dtype: int64

Summary statistics for Traffic Control:

TRAFFIC SIGNAL	46384
----------------	-------

NO CONTROLS	46235
STOP SIGN	9346
FLASHING TRAFFIC SIGNAL	1623
YIELD SIGN	1419
PERSON	200
WARNING SIGN	126
RAILWAY CROSSING DEVICE	35
SCHOOL ZONE SIGN DEVICE	11

Name: Traffic Control, dtype: int64

Summary statistics for Driver Substance Abuse:

NONE DETECTED	84640
DETECTED	20739

Name: Driver Substance Abuse, dtype: int64

Summary statistics for Driver At Fault:

Yes	55157
No	50222

Name: Driver At Fault, dtype: int64

Summary statistics for Injury Severity:

No Injury	84376
Minor Injury	19981
Serious Injury	1022

Name: Injury Severity, dtype: int64

Summary statistics for Driver Distracted By:

NOT DISTRACTED	81402
DISTRACTED	23977

Name: Driver Distracted By, dtype: int64

Summary statistics for Speed Limit:

30-40	73969
15-25	16063
45-55	14587
0-10	688
60-70	72

Name: Speed Limit, dtype: int64

Summary statistics for Location:

(38.953, -77.338)	26
(39.11342767, -77.23648183)	18
(39.11061, -76.98979833)	18
(39.07997592, -77.13826298)	15
(39.045425, -76.99073667)	12
..	
(39.08948038, -77.09106495)	1
(39.089465, -77.13933833)	1

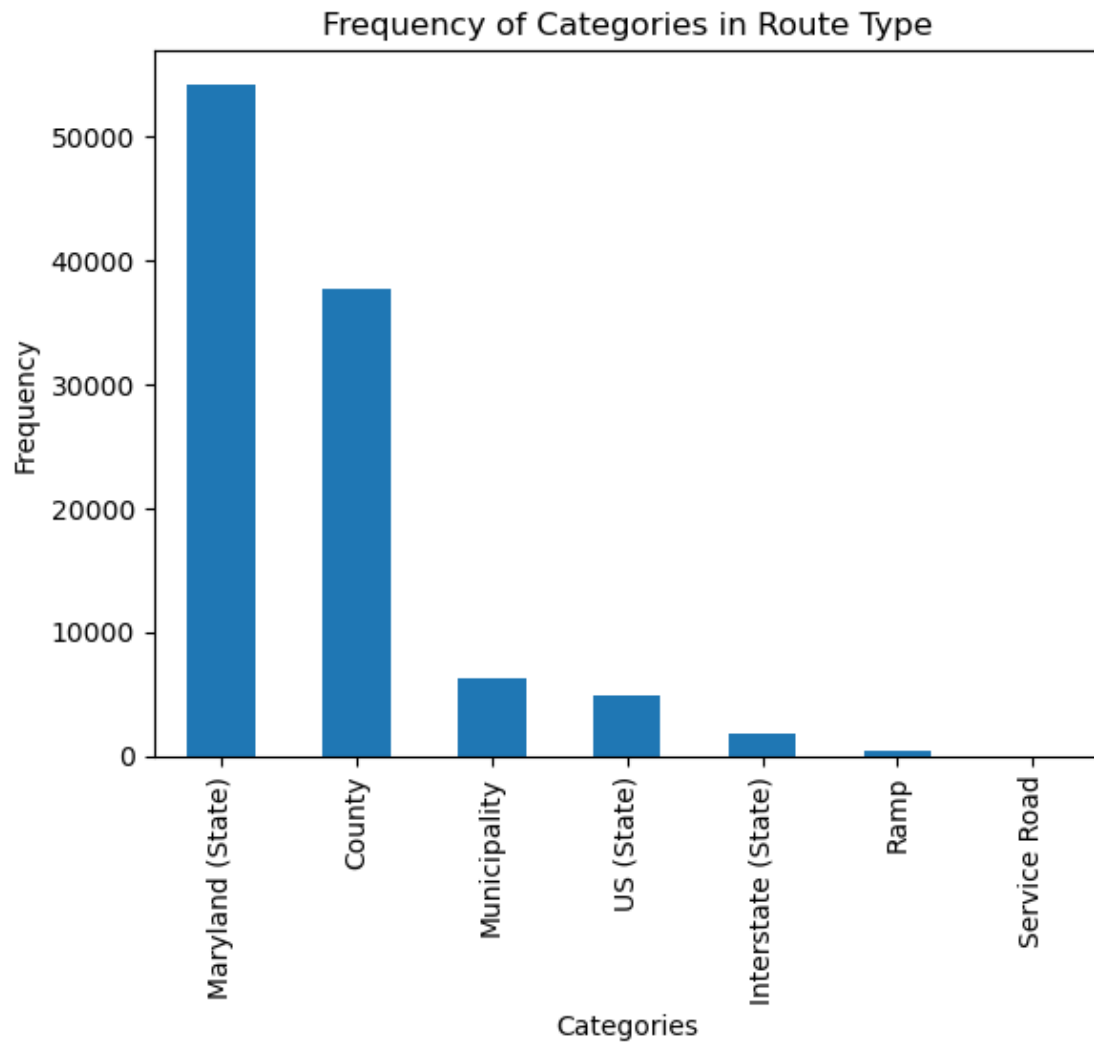
```
(39.15043833, -77.06744833)    1
(39.05174667, -77.06638167)    1
(39.06682667, -77.11454667)    1
Name: Location, Length: 55772, dtype: int64
```


Summary statistics for Day of Week:


```
Friday      17081
Tuesday     16683
Thursday    16517
Wednesday   16425
Monday      15385
Saturday    12657
Sunday      10631
Name: Day of Week, dtype: int64
```


Summary statistics for Time of Day:

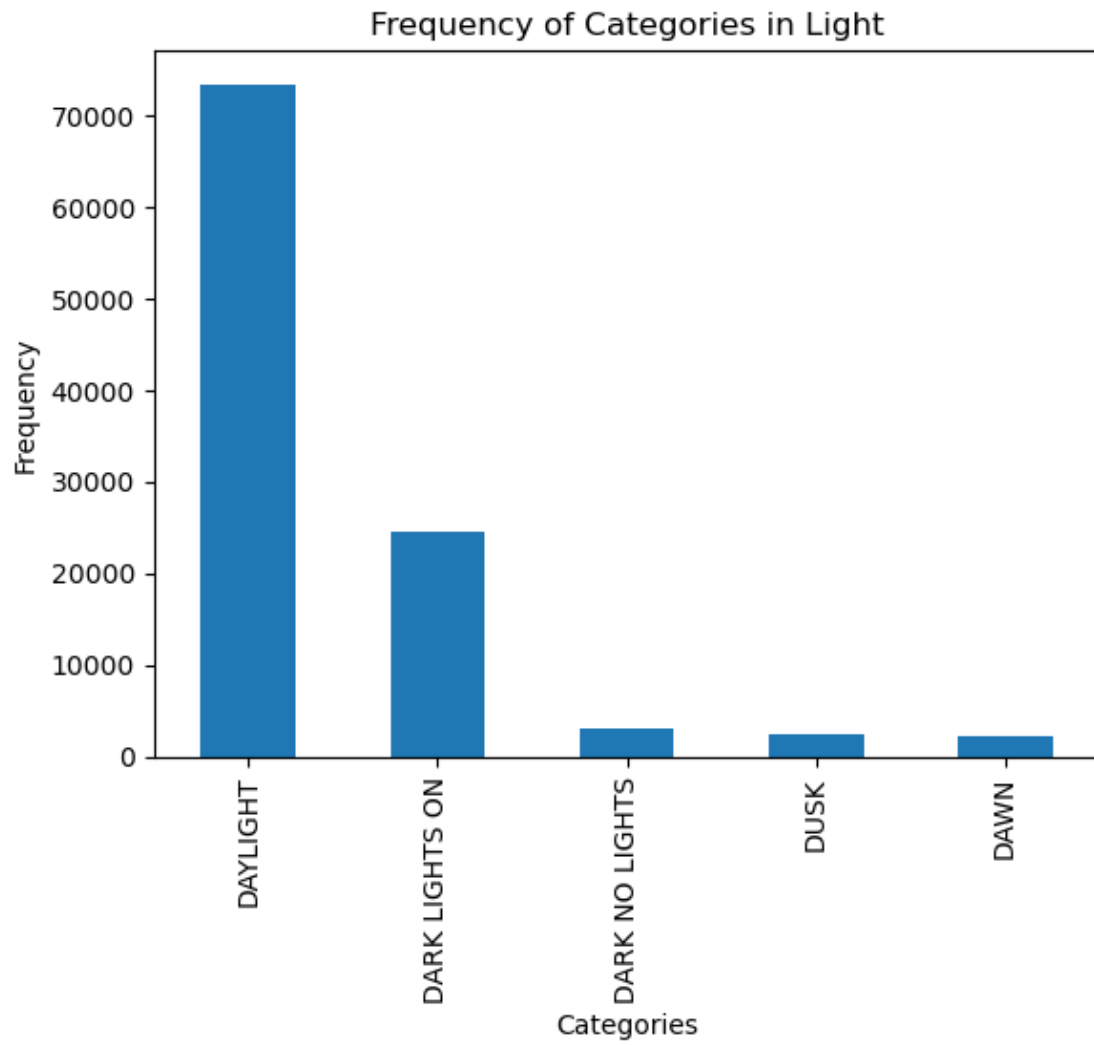
```
afternoon   43517
morning      31297
evening     23983
dawn         6582
Name: Time of Day, dtype: int64
```




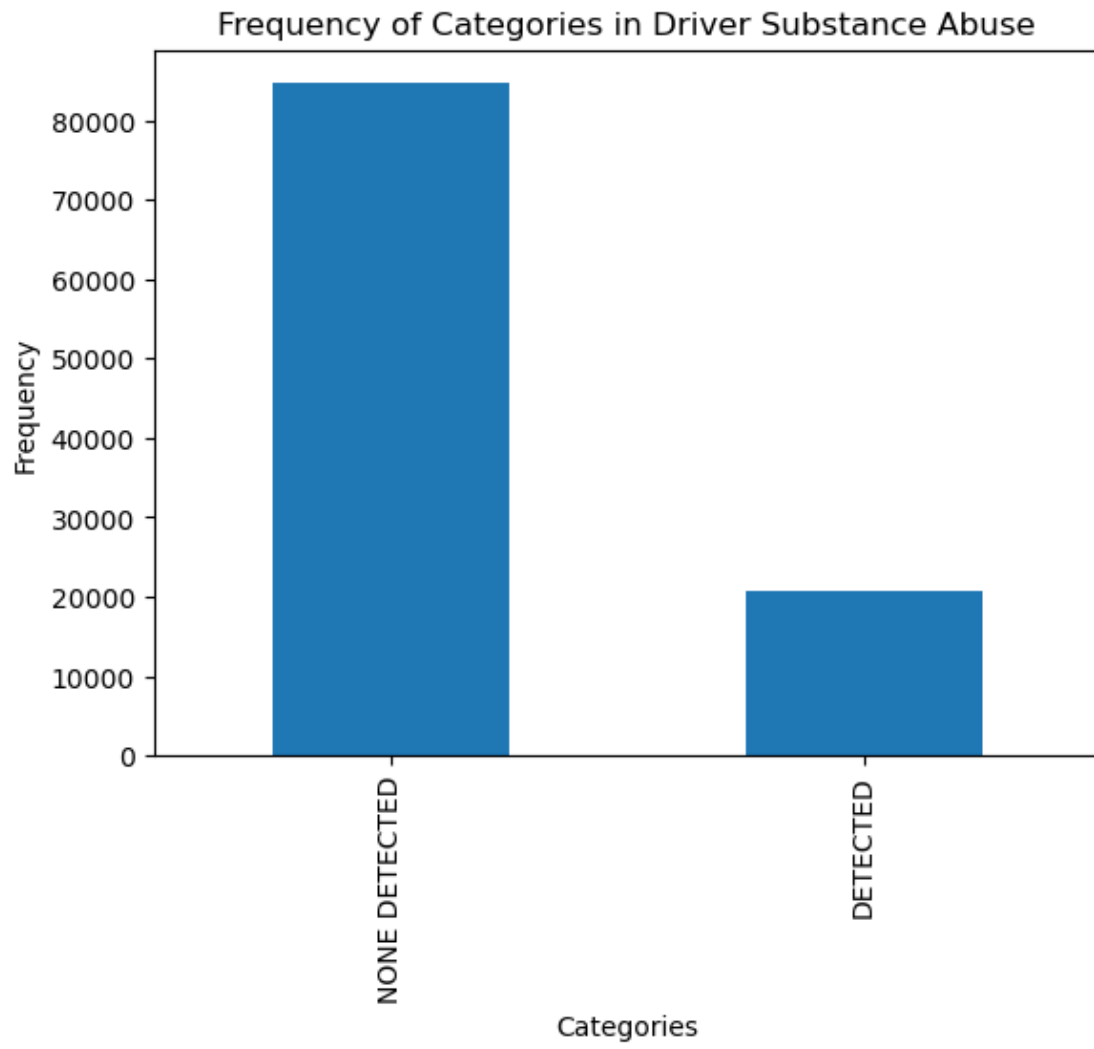


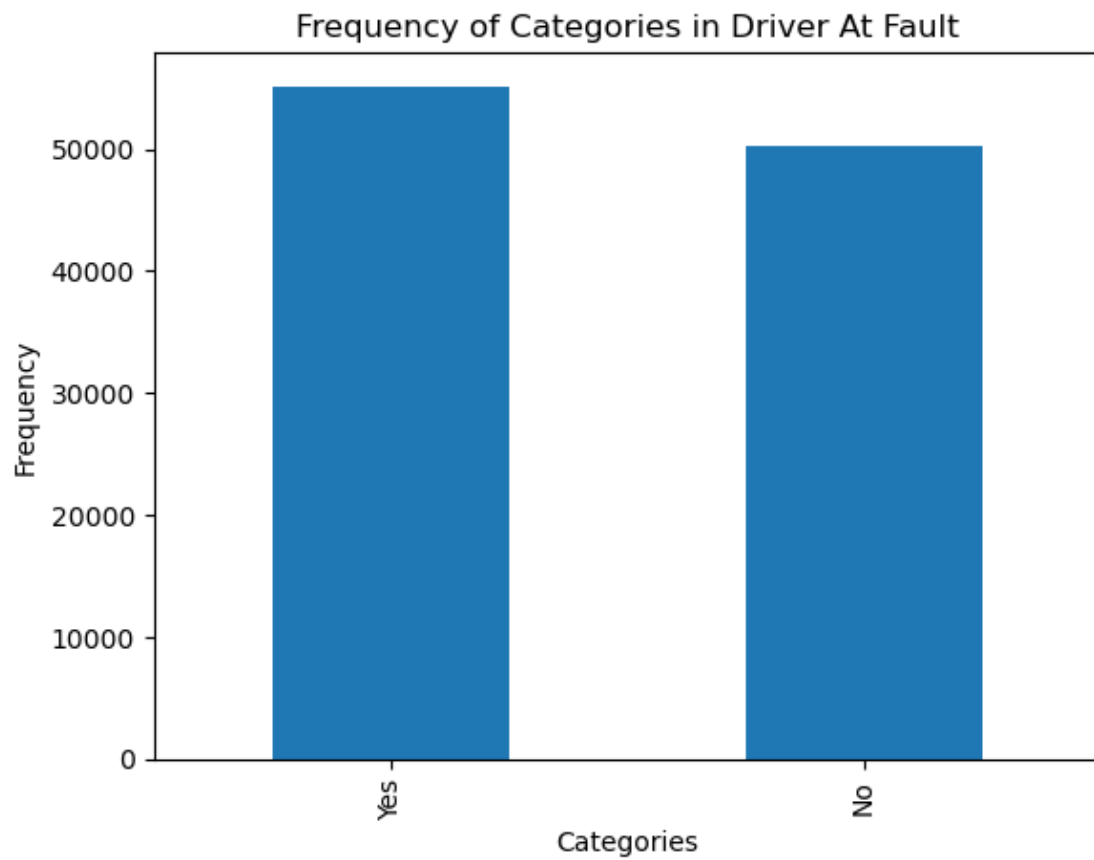


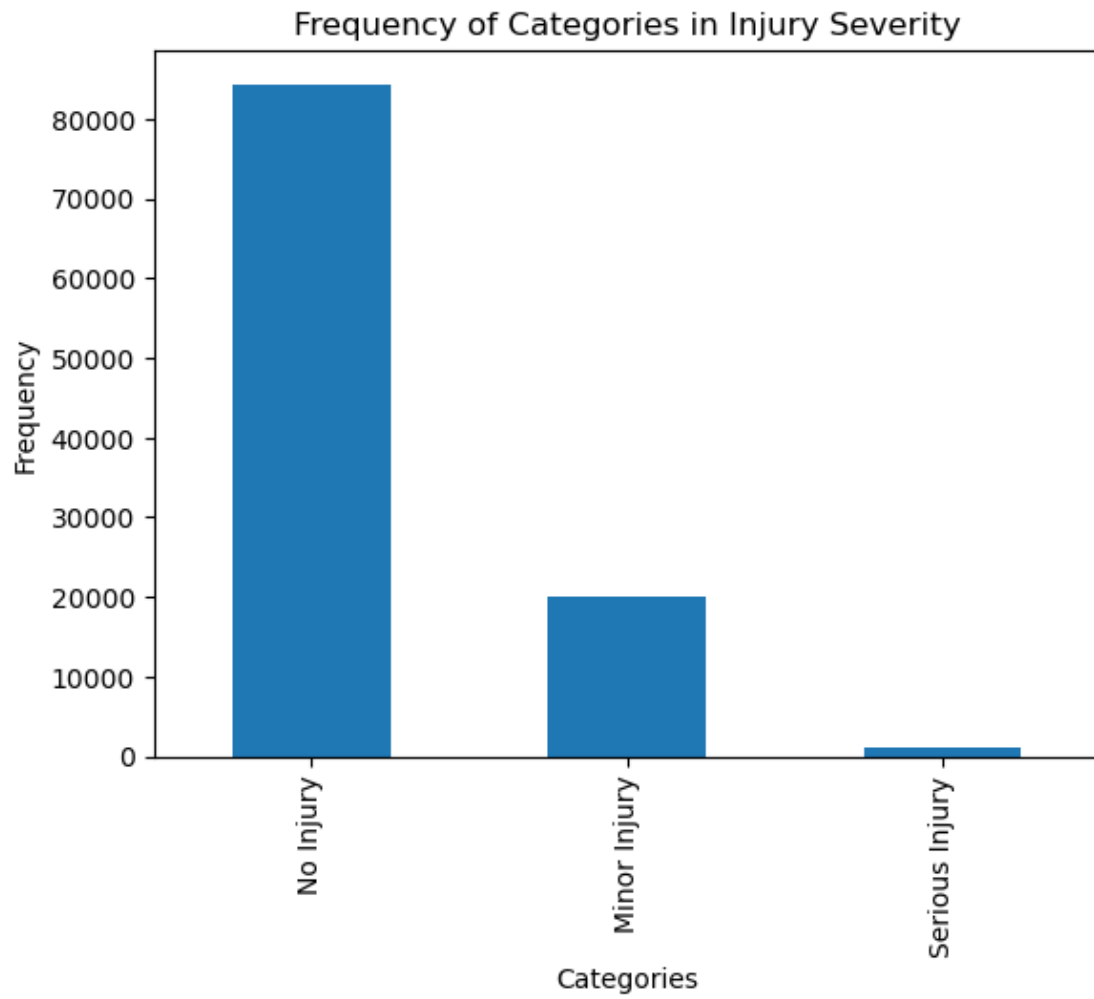


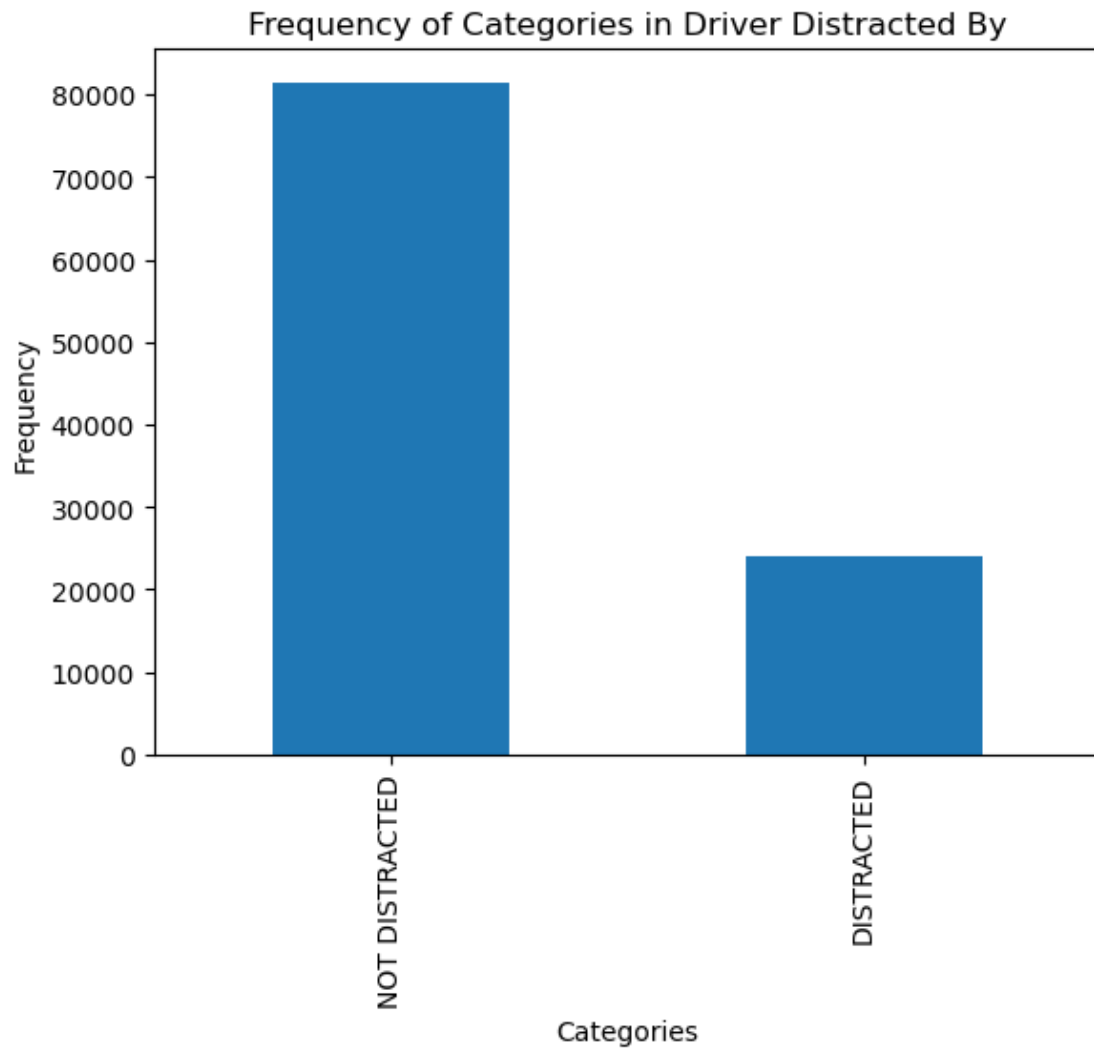


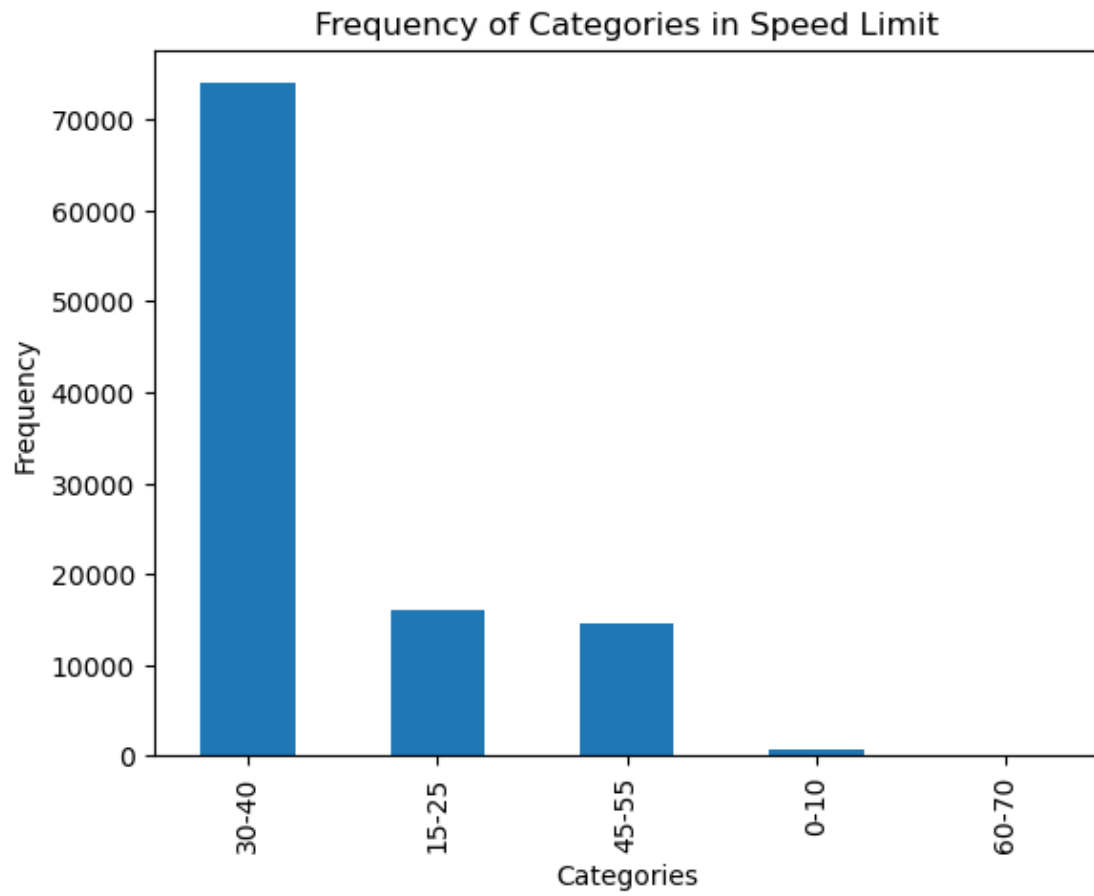





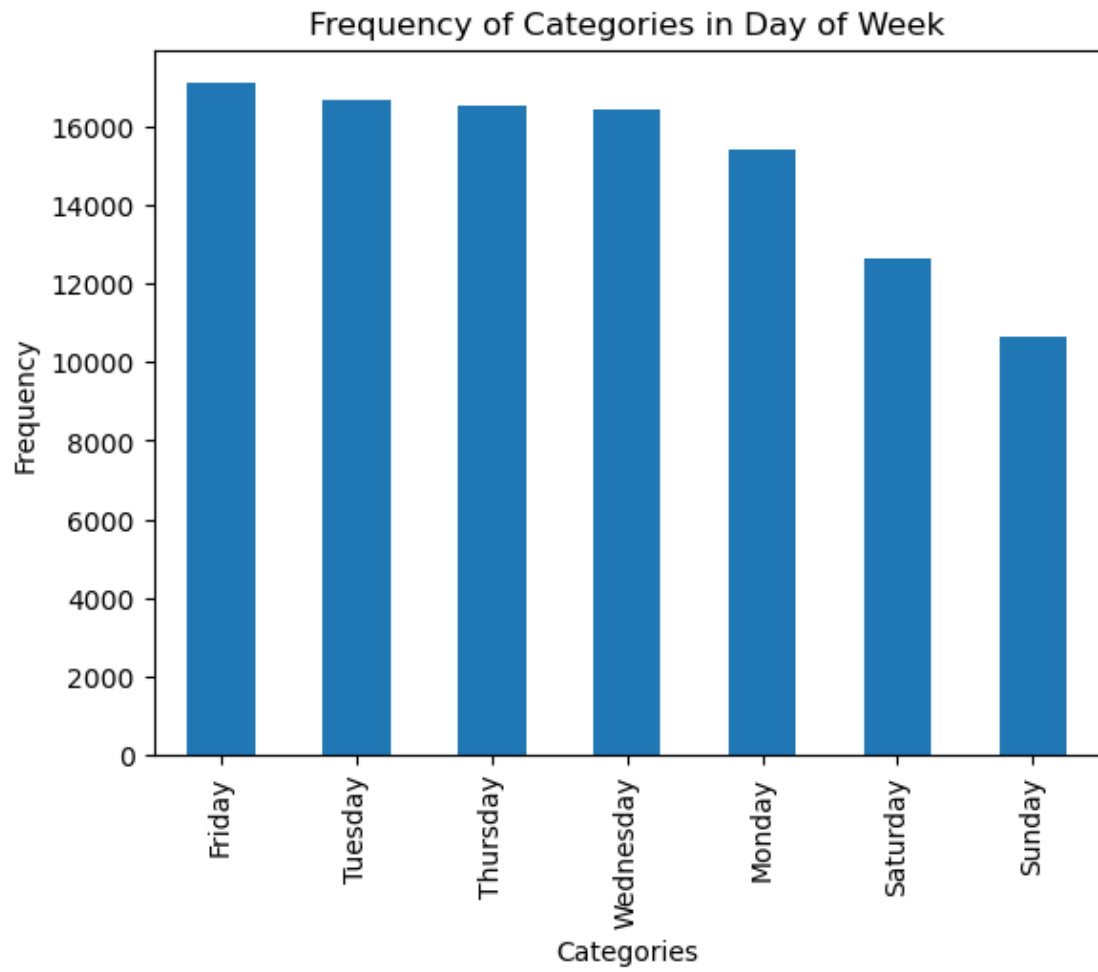


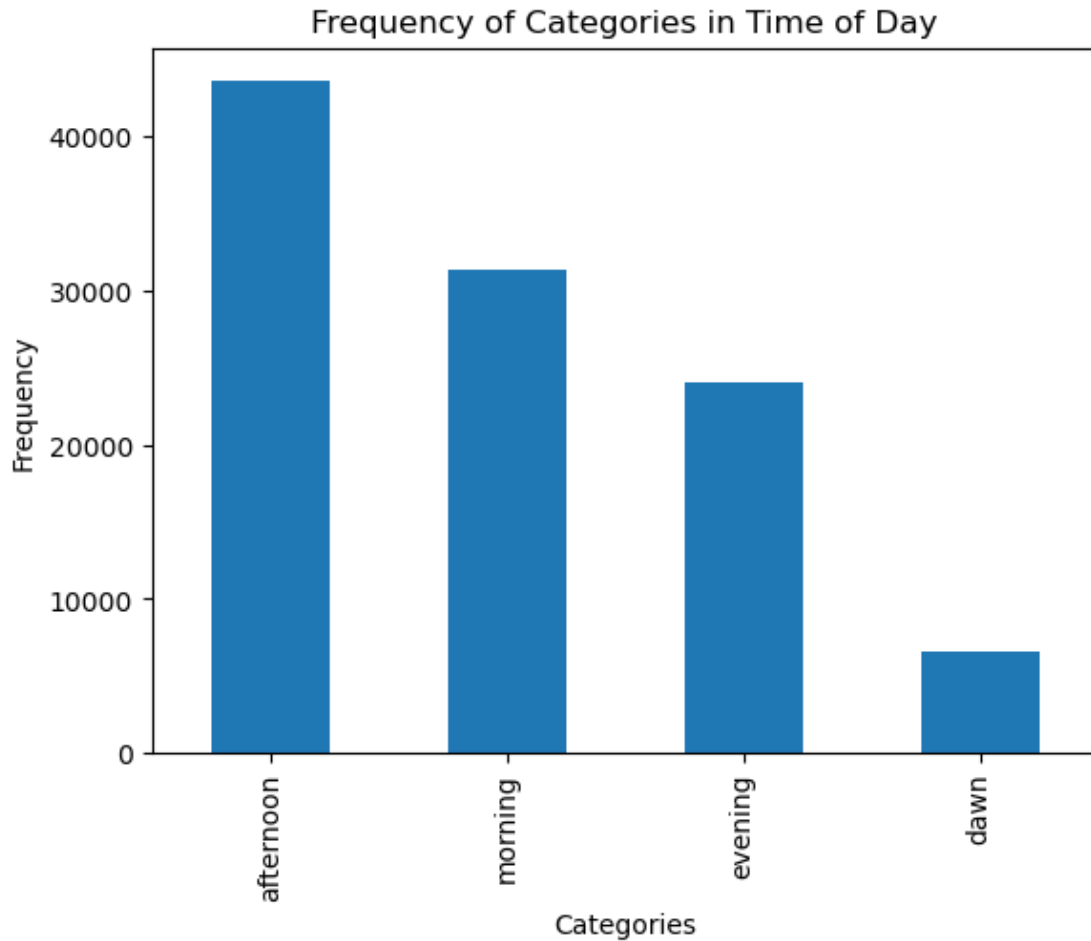













```
[ ]: # spearman correlation matrix and heatmap

import seaborn as sb
import matplotlib.pyplot as plt
import numpy as np

new_df_numeric = new_df.apply(lambda x: x.cat.codes)
correlation_matrix = new_df_numeric.corr(method="spearman", numeric_only=False)
rounded_matrix = np.round(correlation_matrix, decimals=4)
print(rounded_matrix)
plt.figure(figsize=(10,8))
sb.heatmap(rounded_matrix, annot=True, annot_kws={"size": 8}, cmap="coolwarm")
plt.title("Spearman Correlation Matrix")
plt.show()
```

Route Type	Collision Type	Weather	Surface Condition
Light Traffic Control	Driver Substance Abuse	Driver At Fault	Injury Severity
Driver Distracted By	Speed Limit	Location	Day of Week
			Time of Day

Route Type	1.0000	-0.0083	-0.0361	-0.0440
-0.0097	0.0769	0.0031	-0.0416	
0.0000	0.0159	0.1060	0.0053	-0.0057
Collision Type	-0.0083	1.0000	0.0110	0.0077
-0.0181	-0.0439	0.0201	0.0513	
-0.0054	-0.0316	-0.0258	-0.0089	-0.0003
Weather	-0.0361	0.0110	1.0000	0.7511
-0.0111	-0.0037	0.0133	0.0103	
-0.0141	0.0305	0.0034	0.0323	0.0325
Surface Condition	-0.0440	0.0077	0.7511	1.0000
-0.0817	-0.0076	-0.0041	0.0174	
-0.0083	0.0302	0.0125	0.0303	0.0537
Light	-0.0097	-0.0181	-0.0111	-0.0817
1.0000	0.0130	0.1006	-0.0394	
-0.0059	0.0120	-0.0186	0.0001	-0.1083
Traffic Control	0.0769	-0.0439	-0.0037	-0.0076
0.0130	1.0000	-0.0071	-0.0583	
-0.0310	0.0185	0.0885	-0.0066	0.0167
Driver Substance Abuse	0.0031	0.0201	0.0133	-0.0041
0.1006	-0.0071	1.0000	-0.0682	
-0.0265	0.0306	0.0227	-0.0087	0.0085
Driver At Fault	-0.0416	0.0513	0.0103	0.0174
-0.0394	-0.0583	-0.0682	1.0000	
0.1174	-0.4725	-0.0763	0.0078	0.0073
Injury Severity	0.0000	-0.0054	-0.0141	-0.0083
-0.0059	-0.0310	-0.0265	0.1174	
1.0000	-0.0566	-0.0544	-0.0573	-0.0116
Driver Distracted By	0.0159	-0.0316	0.0305	0.0302
0.0120	0.0185	0.0306	-0.4725	
-0.0566	1.0000	0.0472	0.0014	0.0131
Speed Limit	0.1060	-0.0258	0.0034	0.0125
-0.0186	0.0885	0.0227	-0.0763	
-0.0544	0.0472	1.0000	0.1731	0.0125
Location	-0.1024	-0.0089	0.0323	0.0303
0.0001	-0.0066	-0.0087	0.0078	
-0.0573	0.0014	0.1731	1.0000	-0.0021
Day of Week	0.0053	-0.0136	-0.0176	-0.0202
0.0517	-0.0091	0.0172	-0.0083	
-0.0042	0.0015	0.0053	0.0016	0.0222
Time of Day	-0.0057	-0.0003	0.0325	0.0537
-0.1083	0.0167	0.0085	0.0073	
-0.0116	0.0131	0.0125	-0.0021	1.0000



```
[ ]: # very powerful package, however, it is very complicated to use.
      # need to spend more time to understand Cramér's V and dython.
```

```
# import dython as dy
# from dython.nominal import associations
# from dython.nominal import identify_nominal_columns

# associations(new_df)
```

```
[ ]: new_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 105379 entries, 3 to 159356
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---

```

0	Route Type	105379	non-null	category
1	Collision Type	105379	non-null	category
2	Weather	105379	non-null	category
3	Surface Condition	105379	non-null	category
4	Light	105379	non-null	category
5	Traffic Control	105379	non-null	category
6	Driver Substance Abuse	105379	non-null	category
7	Driver At Fault	105379	non-null	category
8	Injury Severity	105379	non-null	category
9	Driver Distracted By	105379	non-null	category
10	Speed Limit	105379	non-null	category
11	Location	105379	non-null	category
12	Day of Week	105379	non-null	category
13	Time of Day	105379	non-null	category

dtypes: category(14)
memory usage: 5.0 MB