

# CIND820\_modelling

June 25, 2023

```
[ ]: # ! pip install imbalanced-learn
import pandas as pd
import sklearn
import numpy as np
```

```
[ ]: df = pd.read_csv("../data/processed/cleaned_data.csv", dtype="category")
df.head()
```

```
[ ]:
Route Type      Collision Type Weather Surface Condition      Light \
0      County              OTHER   CLEAR                DRY      DAYLIGHT
1      County              OTHER  CLOUDY                DRY      DAYLIGHT
2  Municipality  SAME DIR REAR END   CLEAR                DRY      DAWN
3      County      SINGLE VEHICLE  CLOUDY                DRY      DAYLIGHT
4      County      SINGLE VEHICLE   CLEAR                DRY  DARK LIGHTS ON
```

```

Traffic Control Driver Substance Abuse Driver At Fault Injury Severity \
0      NO CONTROLS              DETECTED                Yes      No Injury
1      NO CONTROLS              NONE DETECTED            Yes      Minor Injury
2  TRAFFIC SIGNAL              NONE DETECTED                No      No Injury
3      NO CONTROLS              NONE DETECTED                No      No Injury
4      NO CONTROLS              DETECTED                No      No Injury
```

```

Driver Distracted By Speed Limit Day of Week Time of Day
0      NOT DISTRACTED      15-25      Sunday      afternoon
1      NOT DISTRACTED      15-25      Monday       morning
2      NOT DISTRACTED      30-40     Tuesday       morning
3      NOT DISTRACTED      30-40     Tuesday       morning
4      NOT DISTRACTED      30-40    Thursday        dawn
```

```
[ ]: # Unlike PCA, Random Forest can be utilized for categorical data feature
      ↪ selection in classification.
# Random forest can be used before or after conducting a classification
      ↪ algorithm.

# Before conducting a classification algorithm: Feature importance can help you
      ↪ with feature selection
```

```
# After conducting a classification algorithm: you can examine the feature
↳ importance to gain insights into
# which features had the most influence on the model's predictions. This
↳ retrospective analysis can help you
# understand the key factors driving the classification outcomes and interpret
↳ the model's behavior.
```

```
from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier()
df_numeric = df.apply(lambda x: x.cat.codes)
y_numeric=df_numeric["Injury Severity"]
x_numeric=df_numeric.drop(["Injury Severity"],axis=1)
rf.fit(x_numeric, y_numeric)
```

```
[ ]: RandomForestClassifier()
```

```
[ ]: importance = rf.feature_importances_
indices = np.argsort(importance)[::-1] # Sort indices in descending order
for i, feature_index in enumerate(indices):
    print(f"Feature {i+1}: {x_numeric.columns[feature_index]}")
    ↳ ({importance[feature_index]})"
```

```
Feature 1: Day of Week (0.24150974858501528)
Feature 2: Collision Type (0.1696195607792837)
Feature 3: Route Type (0.09373584935280219)
Feature 4: Time of Day (0.0869449529879933)
Feature 5: Light (0.07894559692671745)
Feature 6: Traffic Control (0.07730874596578606)
Feature 7: Weather (0.07452115068470318)
Feature 8: Speed Limit (0.051921275845526686)
Feature 9: Surface Condition (0.036270814441793454)
Feature 10: Driver Substance Abuse (0.03421606441378519)
Feature 11: Driver At Fault (0.03059462871459191)
Feature 12: Driver Distracted By (0.024411611302001646)
```

```
[ ]: # Due to the nature of traffic accidents,
# the class attribute is significantly imbalanced as you can see from the
↳ output below.
df["Injury Severity"].value_counts()

# This is why balance class attribute is needed in the next code block.
```

```
[ ]: No Injury      84376
Minor Injury      19981
Serious Injury    1022
Name: Injury Severity, dtype: int64
```

```
[ ]: # balance class attribute data
from sklearn.datasets import make_classification

# separate class attribute out before balance
y=df["Injury Severity"]
x=df.drop(["Injury Severity"],axis=1)
nrow=df.shape[0]
ncol=df.shape[1]

class_distribution = df["Injury Severity"].value_counts().to_list()

total= sum(class_distribution)

class_weights = {0: class_distribution[0]/total,
                  1: class_distribution[1]/total,
                  2: class_distribution[2]/total}

# before conducting ROSE or SMOTE, we need to create a synthetic classification
↳ dataset with controlled characteristics.
x, y = make_classification(
    n_samples=nrow, # the number of rows in clean dataset
    n_features=ncol-1, # Total number of features excluding the class attribute
    n_informative=ncol-1, # Number of informative features in your dataset
    n_redundant=0, # Number of redundant features
    n_repeated=0, # Number of repeated features
    n_classes=3, # Number of classes in class attribute
    weights=class_weights, # Class distribution of the target variable
    random_state=42)

print("Generated class distribution:")
print(np.bincount(y))

# if you need to split data into 20% test set and 80% training set.
# x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
↳ random_state=42)
```

Generated class distribution:  
[83894 20148 1337]

```
[ ]: # SMOTE(Synthetic Minority Over-sampling Technique)
# this is a popular technique used to address class imbalance in classification
↳
# it's designed to handle the minority class is underrepresented compared to
↳ the majority class.

from imblearn.over_sampling import SMOTE
```

```

smote = SMOTE(random_state=42)
x_s_resampled, y_s_resampled = smote.fit_resample(x, y)

print("Before SMOTE:")
print(np.bincount(y))

print("After SMOTE:")
print(np.bincount(y_s_resampled))

```

Before SMOTE:  
[83894 20148 1337]  
After SMOTE:  
[83894 83894 83894]

```

[ ]: # cross validation: Stratified K-fold Cross-Validation
# Similar to K-fold, but it ensures that each fold maintains the same class
↪ distribution as the original dataset.

from sklearn.model_selection import StratifiedKFold

k=10
stratified_kfold = StratifiedKFold(n_splits=k, shuffle=True, random_state=42)
# Iterate over the folds
for train_index, test_index in stratified_kfold.split(x_s_resampled,
↪ y_s_resampled):
    # Obtain the training and testing sets for this fold
    x_train, x_test = x_s_resampled[train_index], x_s_resampled[test_index]
    y_train, y_test = y_s_resampled[train_index], y_s_resampled[test_index]

```

```

[ ]: from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score

```

```

[ ]: # for training model, there are 5 algorithms have been selected: logistic
↪ regression, Naive bayes, KNN,
# decision tree and gradient boosting algorithm

# logistic regression

from sklearn.linear_model import LogisticRegression

lr_model=LogisticRegression(
    random_state=42,
    solver="newton-cg",
    warm_start=True).fit(x_train,y_train)
y_pred=lr_model.predict(x_test)
lr_accuracy=accuracy_score(y_test,y_pred)

```

```
f1=f1_score(y_test,y_pred,average="weighted")
print("Accuracy:", lr_accuracy)
print("F1 Score: %.2f" % f1)
```

Accuracy: 0.6612762237762237  
F1 Score: 0.66

```
[ ]: # Naive bayes
```

```
from sklearn.naive_bayes import GaussianNB

nb_model=GaussianNB().fit(x_train,y_train)
y_pred=nb_model.predict(x_test)
nb_accuracy=accuracy_score(y_test,y_pred)
f1=f1_score(y_test,y_pred,average="weighted")
print("Accuracy:", nb_accuracy)
print("F1 Score: %.2f" % f1)

# TD: After the initial result, the Naive Bayes and Logistic Regression
# Seems under - fitting. We will attempt to enhance the performance before
# the final report step.
```

Accuracy: 0.6569055944055944  
F1 Score: 0.66

```
[ ]: # KNN
```

```
from sklearn.neighbors import KNeighborsClassifier

knn_model=KNeighborsClassifier(n_neighbors=500,n_jobs=8).fit(x_train,y_train)
y_pred=knn_model.predict(x_test)
knn_accuracy=accuracy_score(y_test,y_pred)
f1=f1_score(y_test,y_pred,average="weighted")
print("Accuracy:", knn_accuracy)
print("F1 Score: %.2f" % f1)

#TD: After the initial result, we found out the KNN model seems over fitting.
# This issue will be handled before the final report.
```

Accuracy: 0.9034885568976478  
F1 Score: 0.90

```
[ ]: # decision tree
```

```
from sklearn.tree import DecisionTreeClassifier
```

```

ct_model=DecisionTreeClassifier(max_depth=10,random_state=42).
    ↪fit(x_train,y_train)
y_pred=ct_model.predict(x_test)
ct_accuracy=accuracy_score(y_test,y_pred)
f1=f1_score(y_test,y_pred,average="weighted")
print("Accuracy:", ct_accuracy)
print("F1 Score: %.2f" % f1)

# This works perfectly so far. Not over or under fitting.

```

Accuracy: 0.8203671328671329

F1 Score: 0.82

```

[ ]: # gradient boosting algorithm

from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import accuracy_score

gb_model=GradientBoostingClassifier(random_state=42).fit(x_train,y_train)
y_pred=gb_model.predict(x_test)
gb_accuracy=accuracy_score(y_test,y_pred)
f1=f1_score(y_test,y_pred,average="weighted")
print("Accuracy:", gb_accuracy)
print("F1 Score: %.2f" % f1)

# This was not mentioned during the class.
# However, gradient boosting algorithm are used 4 out of 6 journals
# from the literature review section.
# We decided to use this and the performance seems reasonable so far.
# However, this algorithm is quite resource consuming.
# Also, it takes a very long time to finish.

```

Accuracy: 0.8136919898283534

F1 Score: 0.81